

```

import pandas
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score
from sklearn.utils import resample
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.naive_bayes import GaussianNB
from sklearn import svm

df = pd.read_csv(r"C:\Users\nimes\OneDrive\Documents\datasets\
diabetes_dataset.csv")
df.describe()

```

	Age	Pregnancies	BMI	Glucose
BloodPressure \				
count	9538.000000	9538.000000	9538.000000	9538.000000
9538.000000				
mean	53.577584	7.986161	27.052364	106.104183
84.475781				
std	20.764651	4.933469	5.927955	21.918590
14.123480				
min	18.000000	0.000000	15.000000	50.000000
60.000000				
25%	36.000000	4.000000	22.870000	91.000000
74.000000				
50%	53.000000	8.000000	27.050000	106.000000
84.000000				
75%	72.000000	12.000000	31.180000	121.000000
94.000000				
max	89.000000	16.000000	49.660000	207.200000
138.000000				

	HbA1c	LDL	HDL	Triglycerides \
count	9538.000000	9538.000000	9538.000000	9538.000000
mean	4.650661	100.133456	49.953418	151.147746
std	0.476395	29.911910	15.242194	48.951627
min	4.000000	-12.000000	-9.200000	50.000000
25%	4.300000	80.100000	39.700000	117.200000
50%	4.600000	99.900000	50.200000	150.550000
75%	5.000000	120.200000	60.200000	185.100000
max	6.900000	202.200000	107.800000	345.800000

	WaistCircumference	HipCircumference	WHR
FamilyHistory \			
count	9538.000000	9538.000000	9538.000000
9538.000000			
mean	93.951678	103.060621	0.917400

0.302474			
std	15.594468	13.438827	0.140828
0.459354			
min	40.300000	54.800000	0.420000
0.000000			
25%	83.400000	94.000000	0.820000
0.000000			
50%	93.800000	103.200000	0.910000
0.000000			
75%	104.600000	112.100000	1.010000
1.000000			
max	163.000000	156.600000	1.490000
1.000000			

	DietType	Hypertension	MedicationUse	Outcome
count	9538.000000	9538.000000	9538.000000	9538.000000
mean	0.486161	0.001048	0.405012	0.344097
std	0.661139	0.032364	0.490920	0.475098
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	1.000000
max	2.000000	1.000000	1.000000	1.000000

```
target_column = "Outcome"
X = df.drop(columns = target_column)
y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
y_pred = log_reg.predict(X_test)
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
```

Logistic Regression Accuracy: 0.9722222222222222

C:\Users\nimes\PycharmProjects\college\.venv\Lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
rf_reg = RandomForestRegressor()
rf_reg.fit(X_train, y_train)
y_pred = rf_reg.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred))
```

Random Forest Accuracy: 1.0

```
dt_reg = DecisionTreeRegressor()
dt_reg.fit(X_train, y_train)
y_pred = dt_reg.predict(X_test)
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred))
```

Decision Tree Accuracy: 1.0

```
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred = nb.predict(X_test)
print("Naïve Bayes Accuracy:", accuracy_score(y_test, y_pred))
```

Naïve Bayes Accuracy: 0.960167714884696

```
svm_model = svm.SVC(kernel='rbf', C=10, gamma=0.1)
svm_model.fit(X_train, y_train)
y_pred = svm_model.predict(X_test)
print("SVM Accuracy:", f1_score(y_test, y_pred))
```

SVM Accuracy: 0.0

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
print(pd.Series(rf.feature_importances_,
index=df.drop(columns=["Outcome"]).columns).sort_values(ascending=False))
```

FamilyHistory	0.824517
Glucose	0.112114
HbA1c	0.021671
BMI	0.007321
WaistCircumference	0.004873
BloodPressure	0.004463
HipCircumference	0.004433
Age	0.003975
Triglycerides	0.003582
LDL	0.003408
WHR	0.003220
HDL	0.003127
Pregnancies	0.001815
MedicationUse	0.000957
DietType	0.000493

Hypertension	0.000028
dtype: float64	

The observations from this notebook shows that only logistic regression and naive bayes can provide a optimal solution both ranging inbetween 96 to 98 % accuracy, whereas random forest, decision tree and SVM provides the most unreliable outcome.